a Evidence configurations | batch 16, up to 30 epochs

Configuration	Att.	Mean	SD	High	Low + reg.	8 stats	Gate
Attention vector	On	–	–	–	–	–	Learned
+ mean / SD	On	On	On	–	–	–	Learned
+ high-margin vector + 8 statistics	On	On	On	On	–	On	Learned
Full: equal gate	On	On	On	On	On	On	Equal
Full: learned gate	On	On	On	On	On	On	Learned

Inactive vectors are masked; learned gates still normalize over five positions.
The high-margin configuration also enables eight statistics. Main Hybrid uses batch 8.

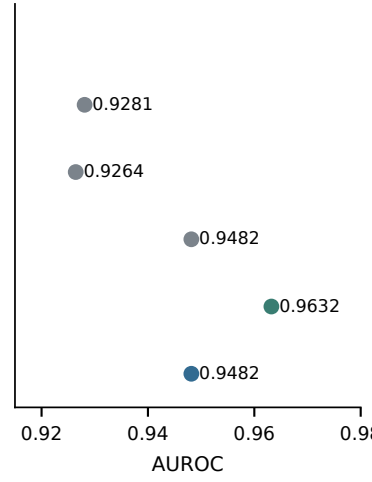

b

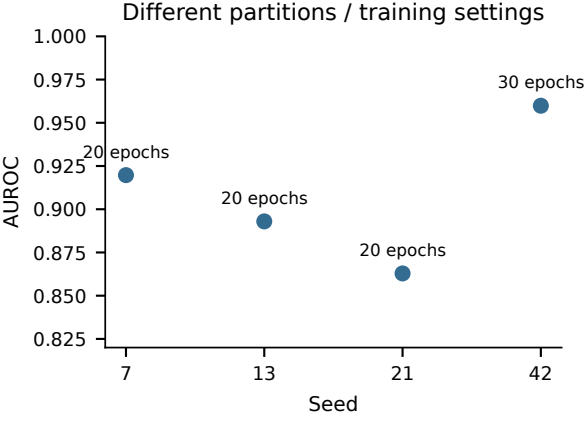

c

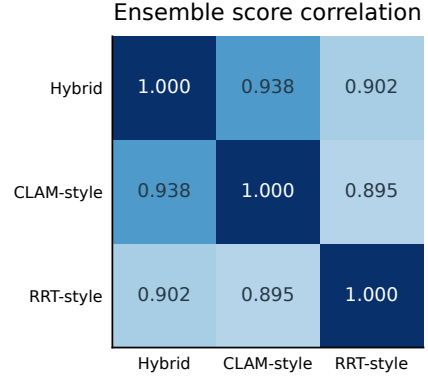

Configuration values and seed values represent individual runs, not matched replicate means.